

✓ Credit Card Fraud Detection

This project aims to detect fraudulent transactions using credit card data. The dataset, containing 284,807 transactions, includes 492 fraud cases, making it highly imbalanced, with fraud accounting for only 0.172% of transactions.

Key Points:

1. Features: The dataset contains only numerical input variables, which are the result of Principal Component Analysis (PCA). Features V1 to V28 represent these components. The 'Time' feature indicates the seconds elapsed since the first transaction, and 'Amount' represents the transaction amount.
2. Target Variable: The 'Class' feature indicates whether a transaction is fraudulent (1) or not (0).
3. Challenge: The dataset is imbalanced, and using accuracy as a performance measure is not reliable. Therefore, we recommend evaluating the model performance using the Area Under the Precision-Recall Curve (AUPRC), which is better suited for imbalanced datasets.

Import Necessary Libraries

```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

# Set seaborn theme for better visual aesthetics
sns.set_theme(style="whitegrid")


# Import necessary libraries
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import (confusion_matrix, accuracy_score,
                             classification_report, precision_recall_curve, auc)
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import RandomizedSearchCV, ShuffleSplit

import warnings
# Suppress warnings for cleaner output
warnings.filterwarnings("ignore")
```

Load DataSet

- Load Data directly through using kaggle API because of data set is little huge which take a long time on google colab

```
from google.colab import files
files.upload()
```

 Choose Files kaggle.json


- **kaggle.json**(application/json) - 66 bytes, last modified: 2/7/2025 - 100% done

Saving kaggle.json to kaggle.json

```
{'kaggle.json': b'{"username": "samikhan25", "key": "e79a0dc30de81afae830d9d1e172ca2c"}'}
```

```
# !pip install kaggle # Install Kaggle API
!mkdir -p ~/.kaggle # Create Kaggle directory
!cp kaggle.json ~/.kaggle/ # Copy your Kaggle API key (Upload kaggle.json first)
!chmod 600 ~/.kaggle/kaggle.json # Set file permissions
```

```
!kaggle datasets download -d mlg-ulb/creditcardfraud
```

 Dataset URL: <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>

License(s): DbCL-1.0

Downloading creditcardfraud.zip to /content

73% 48.0M/66.0M [00:00<00:00, 169MB/s]


100% 66.0M/66.0M [00:00<00:00, 161MB/s]

Unzip dataset


```
from zipfile import ZipFile

dataset="/content/creditcardfraud.zip"

with ZipFile(dataset, 'r') as zip:
    zip.extractall()
    print("The DataSet us extracted")
```


 The DataSet us extracted

```
df = pd.read_csv("/content/creditcard.csv")
df.shape
```



 (284807, 31)

Show Top 5 Rows

```
df.head().T
```



	0	1	2	3	4
Time	0.000000	0.000000	1.000000	1.000000	2.000000
V1	-1.359807	1.191857	-1.358354	-0.966272	-1.158233
V2	-0.072781	0.266151	-1.340163	-0.185226	0.877737
V3	2.536347	0.166480	1.773209	1.792993	1.548718
V4	1.378155	0.448154	0.379780	-0.863291	0.403034
V5	-0.338321	0.060018	-0.503198	-0.010309	-0.407193
V6	0.462388	-0.082361	1.800499	1.247203	0.095921
V7	0.239599	-0.078803	0.791461	0.237609	0.592941
V8	0.098698	0.085102	0.247676	0.377436	-0.270533
V9	0.363787	-0.255425	-1.514654	-1.387024	0.817739
V10	0.090794	-0.166974	0.207643	-0.054952	0.753074
V11	-0.551600	1.612727	0.624501	-0.226487	-0.822843
V12	-0.617801	1.065235	0.066084	0.178228	0.538196
V13	-0.991390	0.489095	0.717293	0.507757	1.345852
V14	-0.311169	-0.143772	-0.165946	-0.287924	-1.119670
V15	1.468177	0.635558	2.345865	-0.631418	0.175121
V16	-0.470401	0.463917	-2.890083	-1.059647	-0.451449
V17	0.207971	-0.114805	1.109969	-0.684093	-0.237033
V18	0.025791	-0.183361	-0.121359	1.965775	-0.038195
V19	0.403993	-0.145783	-2.261857	-1.232622	0.803487
V20	0.251412	-0.069083	0.524980	-0.208038	0.408542
V21	-0.018307	-0.225775	0.247998	-0.108300	-0.009431
V22	0.277838	-0.638672	0.771679	0.005274	0.798278
V23	-0.110474	0.101288	0.909412	-0.190321	-0.137458
V24	0.066928	-0.339846	-0.689281	-1.175575	0.141267
V25	0.128539	0.167170	-0.327642	0.647376	-0.206010
V26	-0.189115	0.125895	-0.139097	-0.221929	0.502292
V27	0.133558	-0.008983	-0.055353	0.062723	0.219422
V28	-0.021053	0.014724	-0.059752	0.061458	0.215153
Amount	149.620000	2.690000	378.660000	123.500000	69.990000
Class	0.000000	0.000000	0.000000	0.000000	0.000000



Dataset information and Statistical Analysis

```
print(df.info())
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Time        284807 non-null  float64
1    V1          284807 non-null  float64
2    V2          284807 non-null  float64
3    V3          284807 non-null  float64
4    V4          284807 non-null  float64
5    V5          284807 non-null  float64
6    V6          284807 non-null  float64
7    V7          284807 non-null  float64
8    V8          284807 non-null  float64
9    V9          284807 non-null  float64
10   V10         284807 non-null  float64
11   V11         284807 non-null  float64
12   V12         284807 non-null  float64
13   V13         284807 non-null  float64
14   V14         284807 non-null  float64
15   V15         284807 non-null  float64
16   V16         284807 non-null  float64
17   V17         284807 non-null  float64
18   V18         284807 non-null  float64
19   V19         284807 non-null  float64
20   V20         284807 non-null  float64
21   V21         284807 non-null  float64
22   V22         284807 non-null  float64
23   V23         284807 non-null  float64
24   V24         284807 non-null  float64
25   V25         284807 non-null  float64
26   V26         284807 non-null  float64
27   V27         284807 non-null  float64
28   V28         284807 non-null  float64
29   Amount      284807 non-null  float64
30   Class       284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
None
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01

8 rows × 31 columns

PREPROCESSING

```
print("before dataset size --> ",df.shape)
print("Missing Values --> ",df.isnull().sum().sum())
print("Duplicated values --> ",df.duplicated().sum())

df.drop_duplicates(inplace=True)
print("\nDrop Sucessfully\n")

print("After drop Missing Values --> ",df.isnull().sum().sum())
print("After drop Duplicated values --> ",df.duplicated().sum())
print("After Drop dataset size --> ",df.shape)
```

```
before dataset size --> (284807, 31)
Missing Values --> 0
Duplicated values --> 1081
```

Drop Sucessfully

```

After drop Missing Values -- >> 0
After drop Duplicated values -- >> 0
After Drop dataset size -- >> (283726, 31)

```

▼ Data Visualization

Class Distribution:

- Highly Imbalanced Dataset (0 vs 1)
- Since Class 0 has 283,253 instances and Class 1 has only 473, it indicates a severe class imbalance.

```

# Set up the figure and subplots
fig, ax = plt.subplots(1, 2, figsize=(12, 4))

# Define colors
colors = ["#ff9999", "#66b3ff"]

# Create a count plot on the first subplot
sns.countplot(x=df["Class"], palette=colors, ax=ax[0])
ax[0].set_title("Class Distribution", fontsize=14, fontweight="bold")
ax[0].set_xlabel("Class", fontsize=12)
ax[0].set_ylabel("Count", fontsize=12)


# Show value counts on top of bars
for p in ax[0].patches:
    ax[0].annotate(f'{p.get_height()}',
                  (p.get_x() + p.get_width() / 2., p.get_height()),
                  ha='center', va='baseline', fontsize=10, fontweight="bold", color='black')

# Create a pie chart on the second subplot
class_counts = df["Class"].value_counts()
ax[1].pie(class_counts, labels=class_counts.index, autopct='%1.1f%%', colors=colors, startangle=90, wedgeprops={'edgecolor': 'black'})

# Set the title for the pie chart
ax[1].set_title("Class Proportion", fontsize=14, fontweight="bold")

# Adjust layout
plt.tight_layout()
plt.show()

```



Numerical Features Distributions

```

# Define the number of features
num_features = len(df.columns)

num_cols = 3 # Fixed number of columns
num_rows = int(np.ceil(num_features / num_cols)) # Calculate required rows

```

```
# Create subplots
fig, ax = plt.subplots(num_rows, num_cols, figsize=(16, num_rows * 3))
ax = ax.flatten()

# Define a color palette
colors = sns.color_palette("husl", num_features)

for i, feature in enumerate(df.columns):
    sns.histplot(df[feature], ax=ax[i], color=colors[i])

    # Set titles and labels
    ax[i].set_title(f"{feature} Distribution", fontsize=10)
    ax[i].set_xlabel(feature, fontsize=8)
    ax[i].set_ylabel("Frequency", fontsize=8)

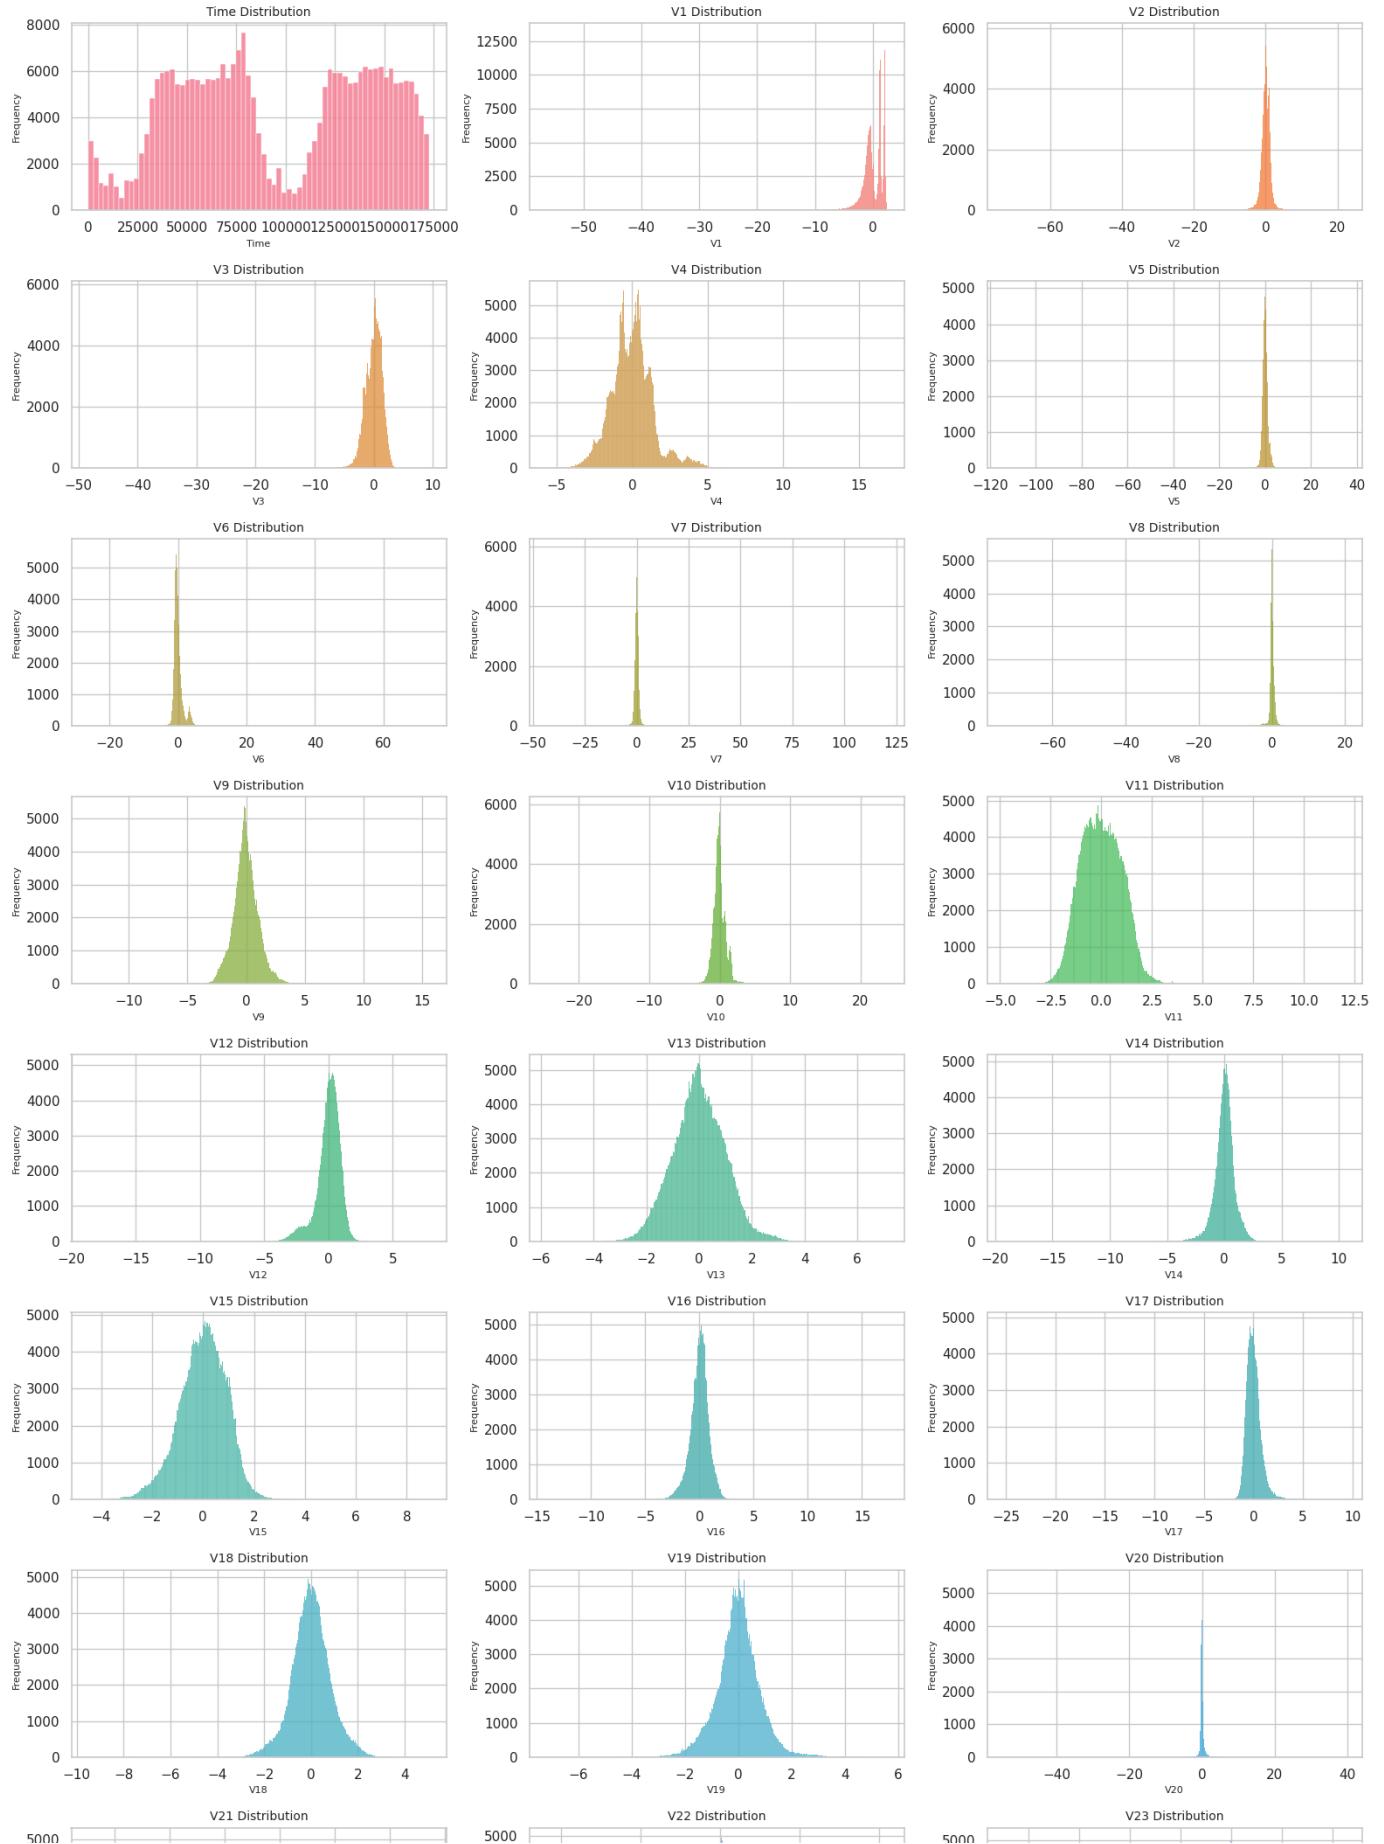
# Hide any unused subplots
for i in range(num_features, len(ax)):
    fig.delaxes(ax[i])

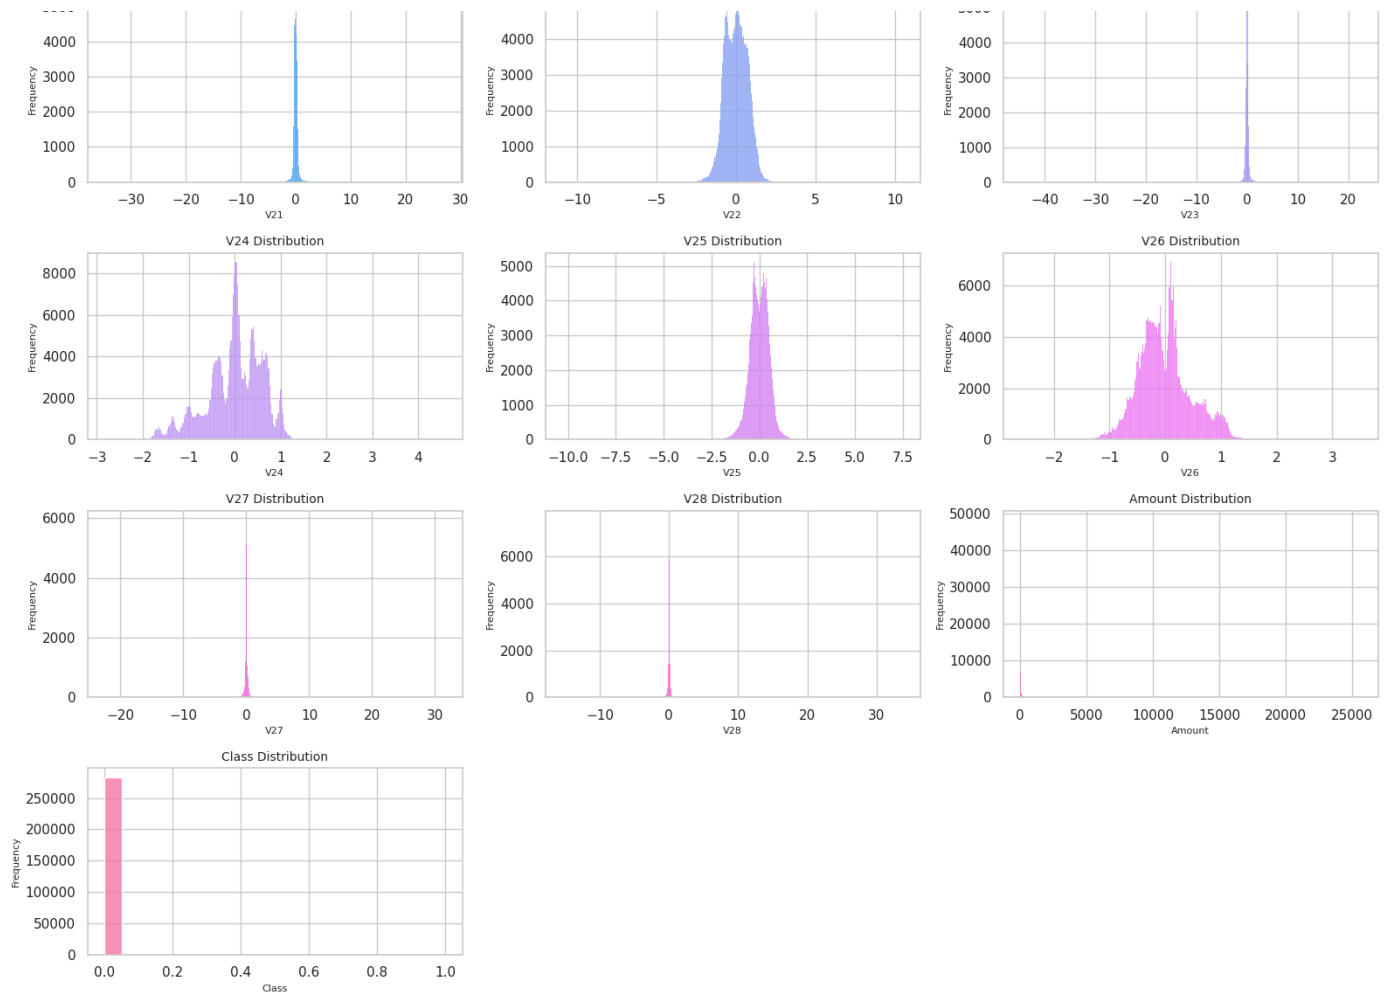
# Set a main title for the entire figure
fig.suptitle("Feature Distributions", fontsize=16, y=1.02)

plt.tight_layout()
plt.show()
```



Feature Distributions






```
num_features = len(df.columns)

num_cols = 3 # Fixed number of columns
num_rows = int(np.ceil(num_features / num_cols)) # Calculate required rows

# Create subplots
fig, ax = plt.subplots(num_rows, num_cols, figsize=(16, num_rows * 3))
ax = ax.flatten()

# Define a color palette
colors = sns.color_palette("husl", num_features)

for i, feature in enumerate(df):
    sns.histplot(df, x=feature, hue=df["Class"].astype(str), ax=ax[i], color=colors[i], multiple='stack')

    # Set titles and labels
    ax[i].set_title(f"{feature} Distribution", fontsize=10)
    ax[i].set_xlabel(feature, fontsize=8)
    ax[i].set_ylabel("Frequency", fontsize=8)

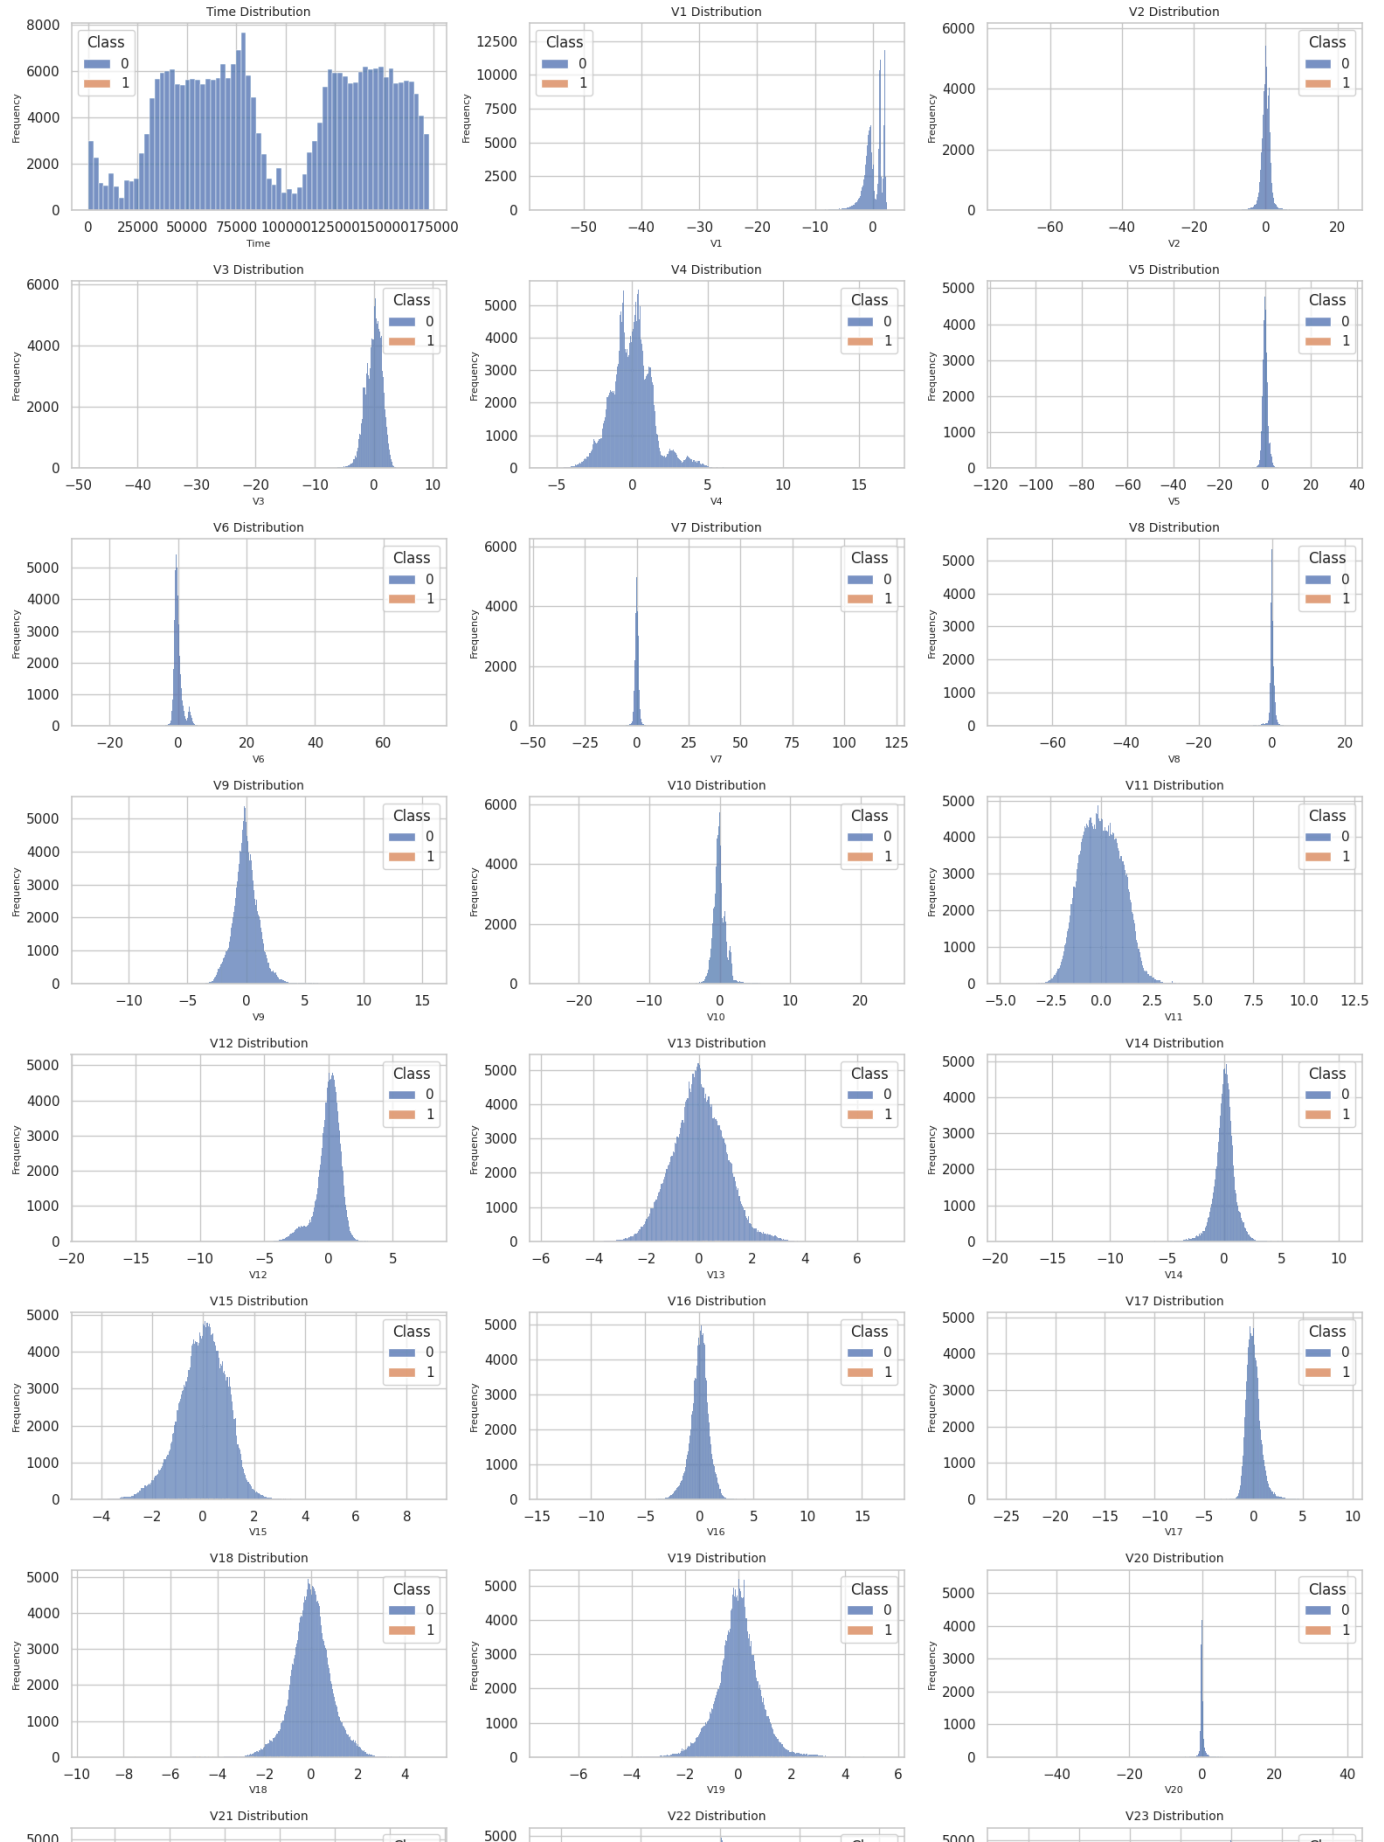
# Hide any unused subplots
for i in range(num_features, len(ax)):
    fig.delaxes(ax[i])

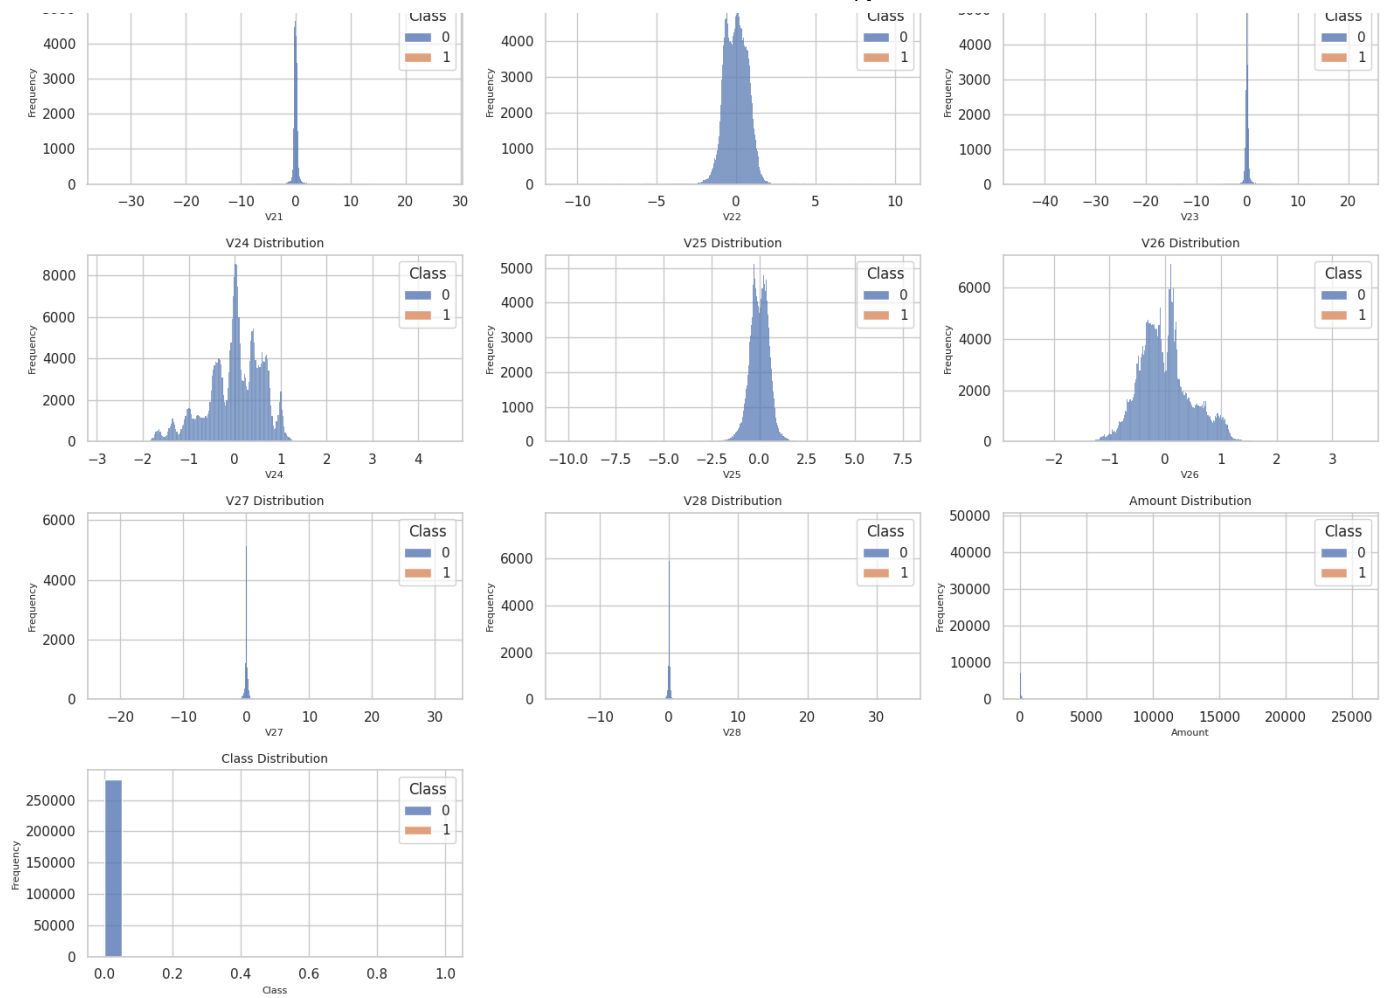
# Set a main title for the entire figure
fig.suptitle("Feature Distributions", fontsize=16, y=1.02)

plt.tight_layout()
plt.show()
```



Feature Distributions





Find Correlation between Features

metrix = df.corr()
metrix


	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22
Time	1.000000	0.117927	-0.010556	-0.422054	-0.105845	0.173223	-0.063279	0.085335	-0.038203	-0.007861	...	0.045913	0.143727
V1	0.117927	1.000000	0.006875	-0.008112	0.002257	-0.007036	0.000413	-0.009173	-0.001168	0.001828	...	0.002818	-0.001436
V2	-0.010556	0.006875	1.000000	0.005278	-0.001495	0.005210	-0.000594	0.007425	0.002899	-0.000274	...	-0.004897	0.001237
V3	-0.422054	-0.008112	0.005278	1.000000	0.002829	-0.006879	-0.001511	-0.011721	-0.001815	-0.003579	...	0.003500	-0.000275
V4	-0.105845	0.002257	-0.001495	0.002829	1.000000	0.001744	-0.000880	0.004657	0.000890	0.002154	...	-0.001034	0.000115
V5	0.173223	-0.007036	0.005210	-0.006879	0.001744	1.000000	-0.000938	-0.008709	0.001430	-0.001213	...	0.001622	-0.000559
V6	-0.063279	0.000413	-0.000594	-0.001511	-0.000880	-0.000938	1.000000	0.000436	0.003036	-0.000734	...	-0.002134	0.001104
V7	0.085335	-0.009173	0.007425	-0.011721	0.004657	-0.008709	0.000436	1.000000	-0.006419	-0.004921	...	0.009010	-0.002280
V8	-0.038203	-0.001168	0.002899	-0.001815	0.000890	0.001430	0.003036	-0.006419	1.000000	0.001038	...	0.018892	-0.006156
V9	-0.007861	0.001828	-0.000274	-0.003579	0.002154	-0.001213	-0.000734	-0.004921	0.001038	1.000000	...	0.000679	0.000785
V10	0.031068	0.000815	0.000620	-0.009632	0.002753	-0.006050	-0.002180	-0.013617	0.000481	-0.012613	...	0.003777	-0.000481
V11	-0.248536	0.001028	-0.000633	0.002339	-0.001223	0.000411	-0.000211	0.002454	0.004688	-0.000217	...	-0.002760	-0.000150
V12	0.125500	-0.001524	0.002266	-0.005900	0.003366	-0.002342	-0.001185	-0.006153	-0.004414	-0.002385	...	0.003285	0.000151
V13	-0.065958	-0.000568	0.000680	0.000113	0.000177	0.000019	0.000397	-0.000170	-0.001381	0.000745	...	0.000522	0.000016
V14	-0.100316	-0.002663	0.002711	-0.003027	0.002801	-0.001000	0.000184	-0.003816	-0.008387	0.001981	...	0.005633	-0.001906
V15	-0.184392	-0.000602	0.001538	-0.001230	0.000572	-0.001171	-0.000470	-0.001394	0.001044	-0.000283	...	-0.000271	-0.001197
V16	0.011286	-0.003345	0.004013	-0.004430	0.003346	-0.002373	0.000122	-0.005944	-0.004376	-0.000086	...	0.004326	-0.000820
V17	-0.073819	-0.003491	0.003244	-0.008159	0.003655	-0.004466	-0.001716	-0.008794	-0.005576	-0.002318	...	0.003560	-0.000162
V18	0.090305	-0.003535	0.002477	-0.003495	0.002325	-0.002685	0.000541	-0.004279	-0.001323	-0.000373	...	0.001629	-0.000533
V19	0.029537	0.000919	-0.000358	-0.000016	-0.000560	0.000436	0.000106	0.000846	-0.000626	0.000247	...	0.000244	0.001342
V20	-0.051022	-0.001393	-0.001287	-0.002269	0.000318	-0.001185	-0.000181	-0.001192	0.000271	-0.001838	...	0.005372	-0.001617
V21	0.045913	0.002818	-0.004897	0.003500	-0.001034	0.001622	-0.002134	0.009010	0.018892	0.000679	...	1.000000	0.009645
V22	0.143727	-0.001436	0.001237	-0.000275	0.000115	-0.000559	0.001104	-0.002280	-0.006156	0.000785	...	0.009645	1.000000
V23	0.051474	-0.001330	-0.003855	0.000449	0.000732	0.001183	-0.000755	0.003303	0.004994	0.000677	...	-0.006391	0.001929
V24	-0.015954	-0.000723	0.000701	-0.000072	-0.000120	0.000198	0.001202	-0.000384	0.000113	-0.000103	...	0.001210	-0.000031
V25	-0.233262	-0.000222	-0.001569	0.000425	0.000162	0.000069	0.000697	-0.000072	0.000011	-0.000275	...	-0.000872	0.000197
V26	-0.041818	-0.000684	0.000253	-0.000094	0.000777	0.000390	-0.000028	0.000624	-0.001407	0.001253	...	-0.000874	-0.001495
V27	-0.005171	-0.015706	0.007555	-0.007051	0.001322	-0.005798	0.000289	-0.004537	0.000613	0.008221	...	-0.005216	0.003037
V28	-0.009305	-0.004861	0.001611	-0.000134	0.000231	-0.000820	0.000925	0.001657	-0.000099	0.005591	...	-0.004436	0.001392
Amount	-0.010559	-0.230105	-0.533428	-0.212410	0.099514	-0.387685	0.216389	0.400408	-0.104662	-0.044123	...	0.108058	-0.064965
Class	-0.012359	-0.094486	0.084624	-0.182322	0.129326	-0.087812	-0.043915	-0.172347	0.033068	-0.094021	...	0.026357	0.004887

31 rows × 31 columns

```
# Plot heatmap
plt.figure(figsize=(14, 12)) # Set figure size
sns.heatmap(metrix, annot=True, cmap="Blues", fmt=".2f", linewidths=0.5,
            linecolor="black", annot_kws={"size": 6})

# title and labels for correlation matrix
plt.title("Feature Correlation Heatmap")
plt.xlabel("Features")
plt.ylabel("Features")

plt.show()
```



Feature Engineering

1. Define X and y: Separate the features and target variable.
2. Split the Data: Split the dataset into training and testing sets.
3. Scale the Dataset: Normalize the feature values for improved model performance.

```
# Define features (X) and target variable (y)
X = df.drop(columns=["Class"]) # Dropping the target column
y = df["Class"] # Target variable
```

```
# Splitting dataset (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
# Print dataset sizes with better formatting
print(f"X_train Size: {X_train.shape}")
print(f"y_train Size: {y_train.shape}")
print(f"X_test Size: {X_test.shape}")
```

```
print(f"y_test Size: {y_test.shape}")
```

```
↗ X_train Size: (226980, 30)
  y_train Size: (226980,)
  X_test Size: (56746, 30)
  y_test Size: (56746,)
```

```
# Initialize the scaler
scaler = StandardScaler()

# Fit on training data and transform both sets
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test) # Only transform test data

# Print confirmation
print("Feature scaling applied both X_train,X_test")
```

```
↗ Feature scaling applied both X_train,X_test
```

✓ Model Trainings

1. LogisticRegression
2. Decission Tree
3. RandomForest
4. Gradien Boosting
5. XGBoost
6. ANN(Artificial Neural Network)

Function include:

- Model Training
- Evaluation
- Performance Metrics Visualization

✓ Resampling

```
# Initialize SMOTE
smote = SMOTE(random_state=42)

# Resample the training data
X_train_resample, y_train_resample = smote.fit_resample(X_train, y_train)

# Check the shape of the resampled data
print(f"Original X_train shape: {X_train.shape}")
print(f"Original y_train shape: {y_train.shape}")
print(f"Resampled X_train shape: {X_train_resample.shape}")
print(f"Resampled y_train shape: {y_train_resample.shape}")
```

```
↗ Original X_train shape: (226980, 30)
  Original y_train shape: (226980,)
  Resampled X_train shape: (453204, 30)
  Resampled y_train shape: (453204,)
```

✓ Model Training on Resample Dataset

```
def model_train_evaluation_resampling(model):

    model.fit(X_train_resample, y_train_resample)

    # Make predictions on both training and test data
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)
```

```
# Evaluate performance
train_acc = accuracy_score(y_train, y_train_pred)
test_acc = accuracy_score(y_test, y_test_pred)

# Print accuracy scores
print(f"Training Accuracy: {train_acc:.4f}")
print(f"Testing Accuracy: {test_acc:.4f}")

# Generate classification reports
print("\nClassification Report (Training Data):\n", classification_report(y_train, y_train_pred))
print("\nClassification Report (Testing Data):\n", classification_report(y_test, y_test_pred))

# Generate confusion matrices
train_cm = confusion_matrix(y_train, y_train_pred)
test_cm = confusion_matrix(y_test, y_test_pred)

# Plot confusion matrix heatmaps
fig, ax = plt.subplots(1, 2, figsize=(12, 5))

# Training Confusion Matrix
sns.heatmap(train_cm, annot=True, fmt="d", cmap="Blues", ax=ax[0])
ax[0].set_title("Confusion Matrix (Training)")
ax[0].set_xlabel("Predicted Label")
ax[0].set_ylabel("True Label")

# Testing Confusion Matrix
sns.heatmap(test_cm, annot=True, fmt="d", cmap="Reds", ax=ax[1])
ax[1].set_title("Confusion Matrix (Testing)")
ax[1].set_xlabel("Predicted Label")
ax[1].set_ylabel("True Label")

plt.tight_layout()
plt.show()
```

Model --> 1 -- >> LogisticRegression

```
log_model_2 = LogisticRegression(class_weight='balanced' , random_state=42)
model_train_evaluation_resampling(log_model_2)
```

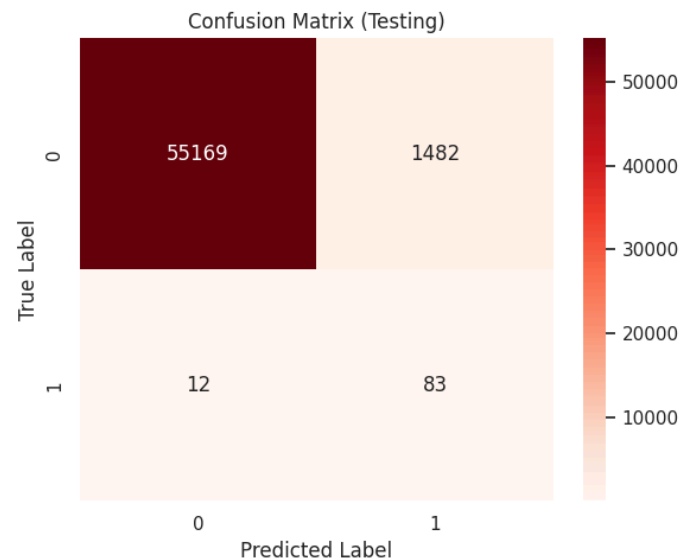

↻ Training Accuracy: 0.9743
Testing Accuracy: 0.9737

Classification Report (Training Data):

	precision	recall	f1-score	support
0	1.00	0.97	0.99	226602
1	0.06	0.93	0.11	378
accuracy			0.97	226980
macro avg	0.53	0.95	0.55	226980
weighted avg	1.00	0.97	0.99	226980

Classification Report (Testing Data):

	precision	recall	f1-score	support
0	1.00	0.97	0.99	56651
1	0.05	0.87	0.10	95
accuracy			0.97	56746
macro avg	0.53	0.92	0.54	56746
weighted avg	1.00	0.97	0.99	56746



Model --> 2 -->>Decision Tree

```
dt_model_2 = DecisionTreeClassifier()
model_train_evaluation_resampling(dt_model_2)
```


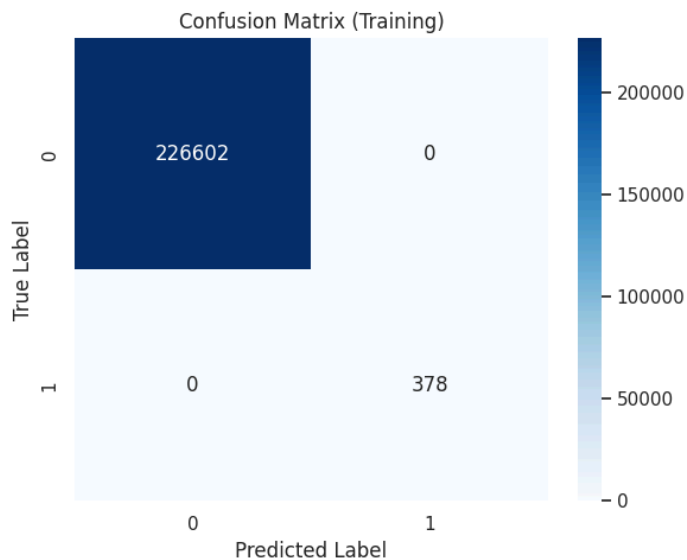

↻ Training Accuracy: 1.0000
Testing Accuracy: 0.9974

Classification Report (Training Data):

	precision	recall	f1-score	support
0	1.00	1.00	1.00	226602
1	1.00	1.00	1.00	378
accuracy			1.00	226980
macro avg	1.00	1.00	1.00	226980
weighted avg	1.00	1.00	1.00	226980

Classification Report (Testing Data):

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56651
1	0.35	0.64	0.45	95
accuracy			1.00	56746
macro avg	0.67	0.82	0.72	56746
weighted avg	1.00	1.00	1.00	56746



Model --> 3 -->> Random Forest

```
rf_model_2 = RandomForestClassifier(n_estimators=500,class_weight='balanced',random_state=42,criterion='gini',n_jobs=-1)
model_train_evaluation_resampling(rf_model_2)
```